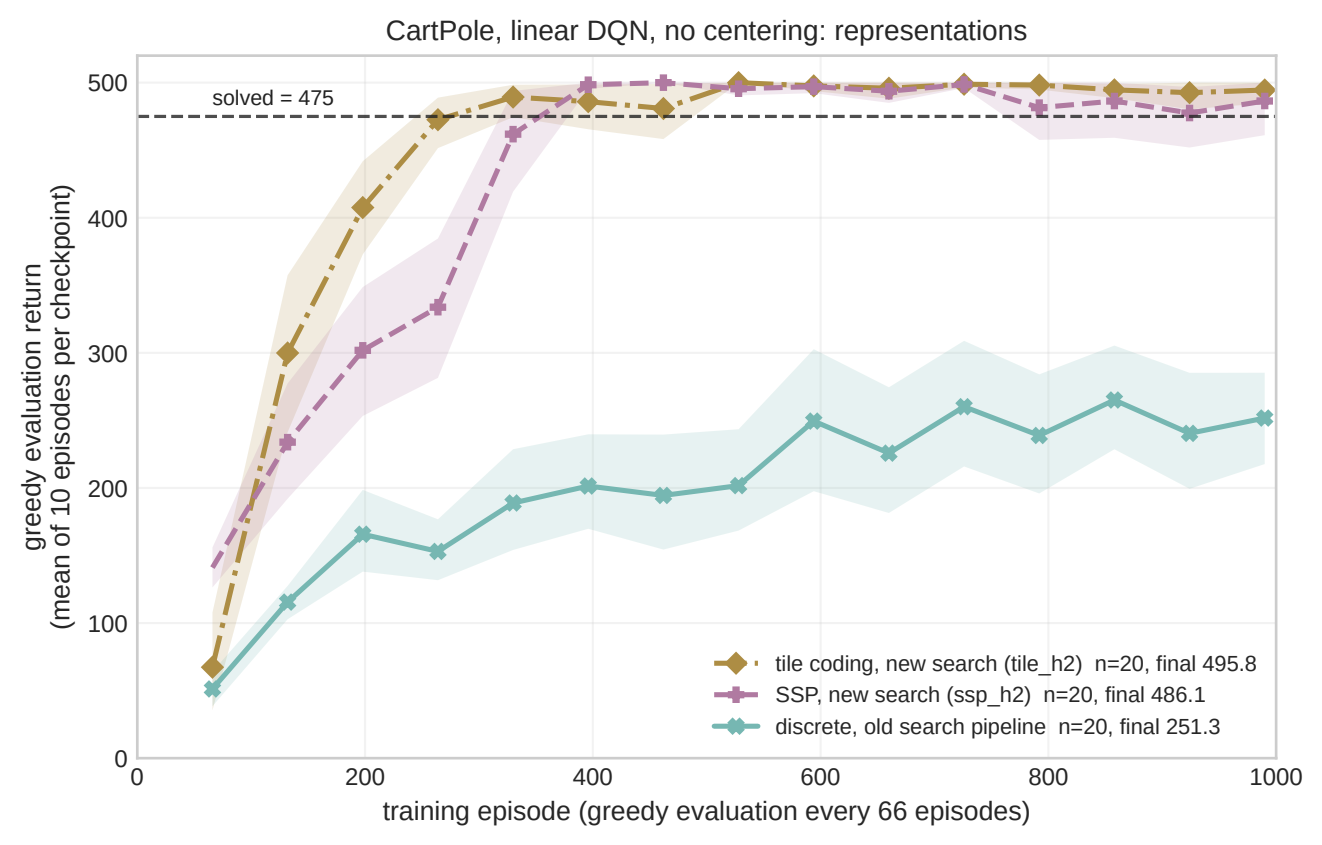
CartPole, linear DQN, no centering: representations
solved = 475
greedy evaluation return
(mean of 10 episodes per checkpoint)
training episode (greedy evaluation every 66 episodes)
0
100
200
300
400
500
0
200
400
600
800
1000
tile coding, new search (tile_h2) n=20, final 495.8
SSP, new search (ssp_h2) n=20, final 486.1
discrete, old search pipeline n=20, final 251.3