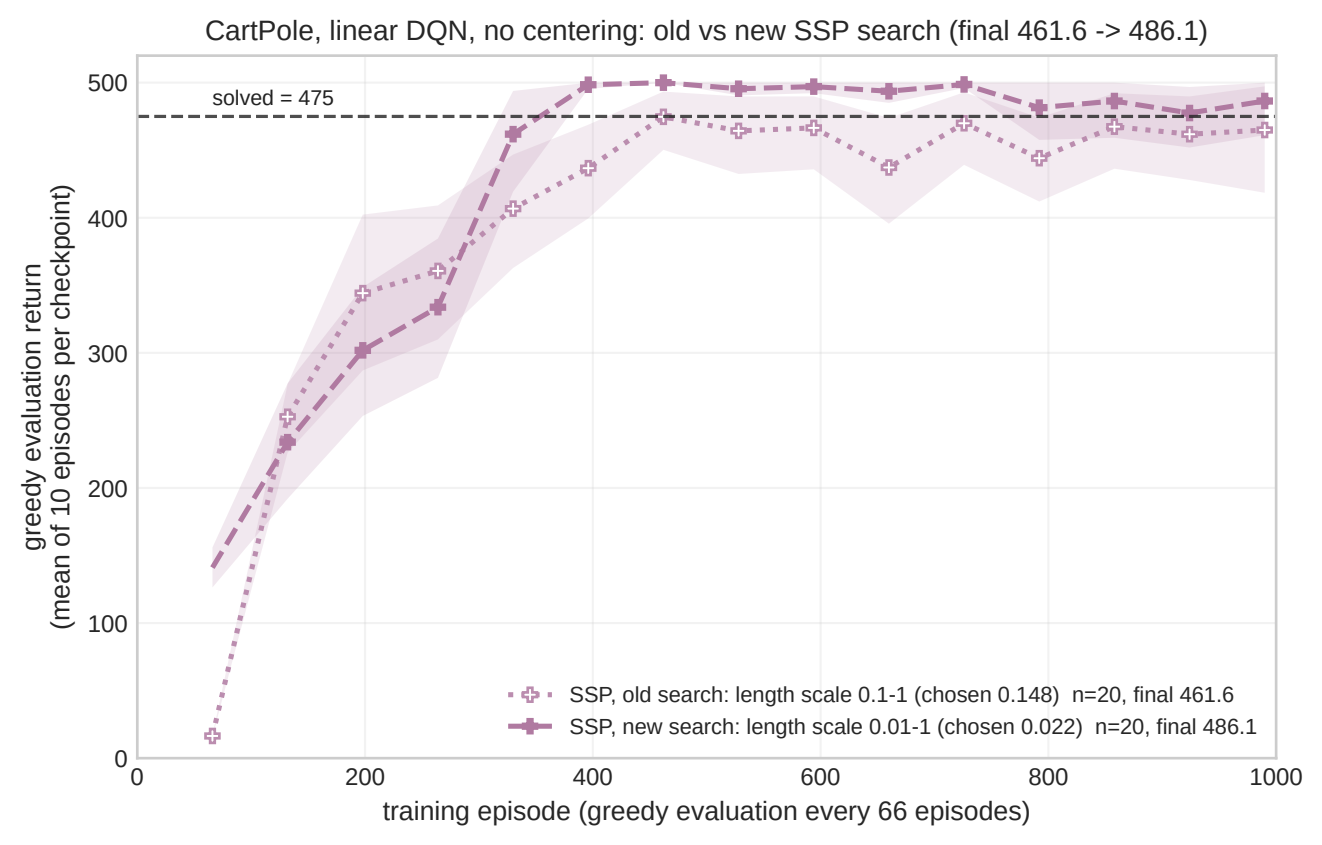
CartPole, linear DQN, no centering: old vs new SSP search (final 461.6 -> 486.1)
solved = 475
greedy evaluation return
(mean of 10 episodes per checkpoint)
training episode (greedy evaluation every 66 episodes)
SSP, old search: length scale 0.1-1 (chosen 0.148) n=20, final 461.6
SSP, new search: length scale 0.01-1 (chosen 0.022) n=20, final 486.1
0
100
200
300
400
500
0
200
400
600
800
1000