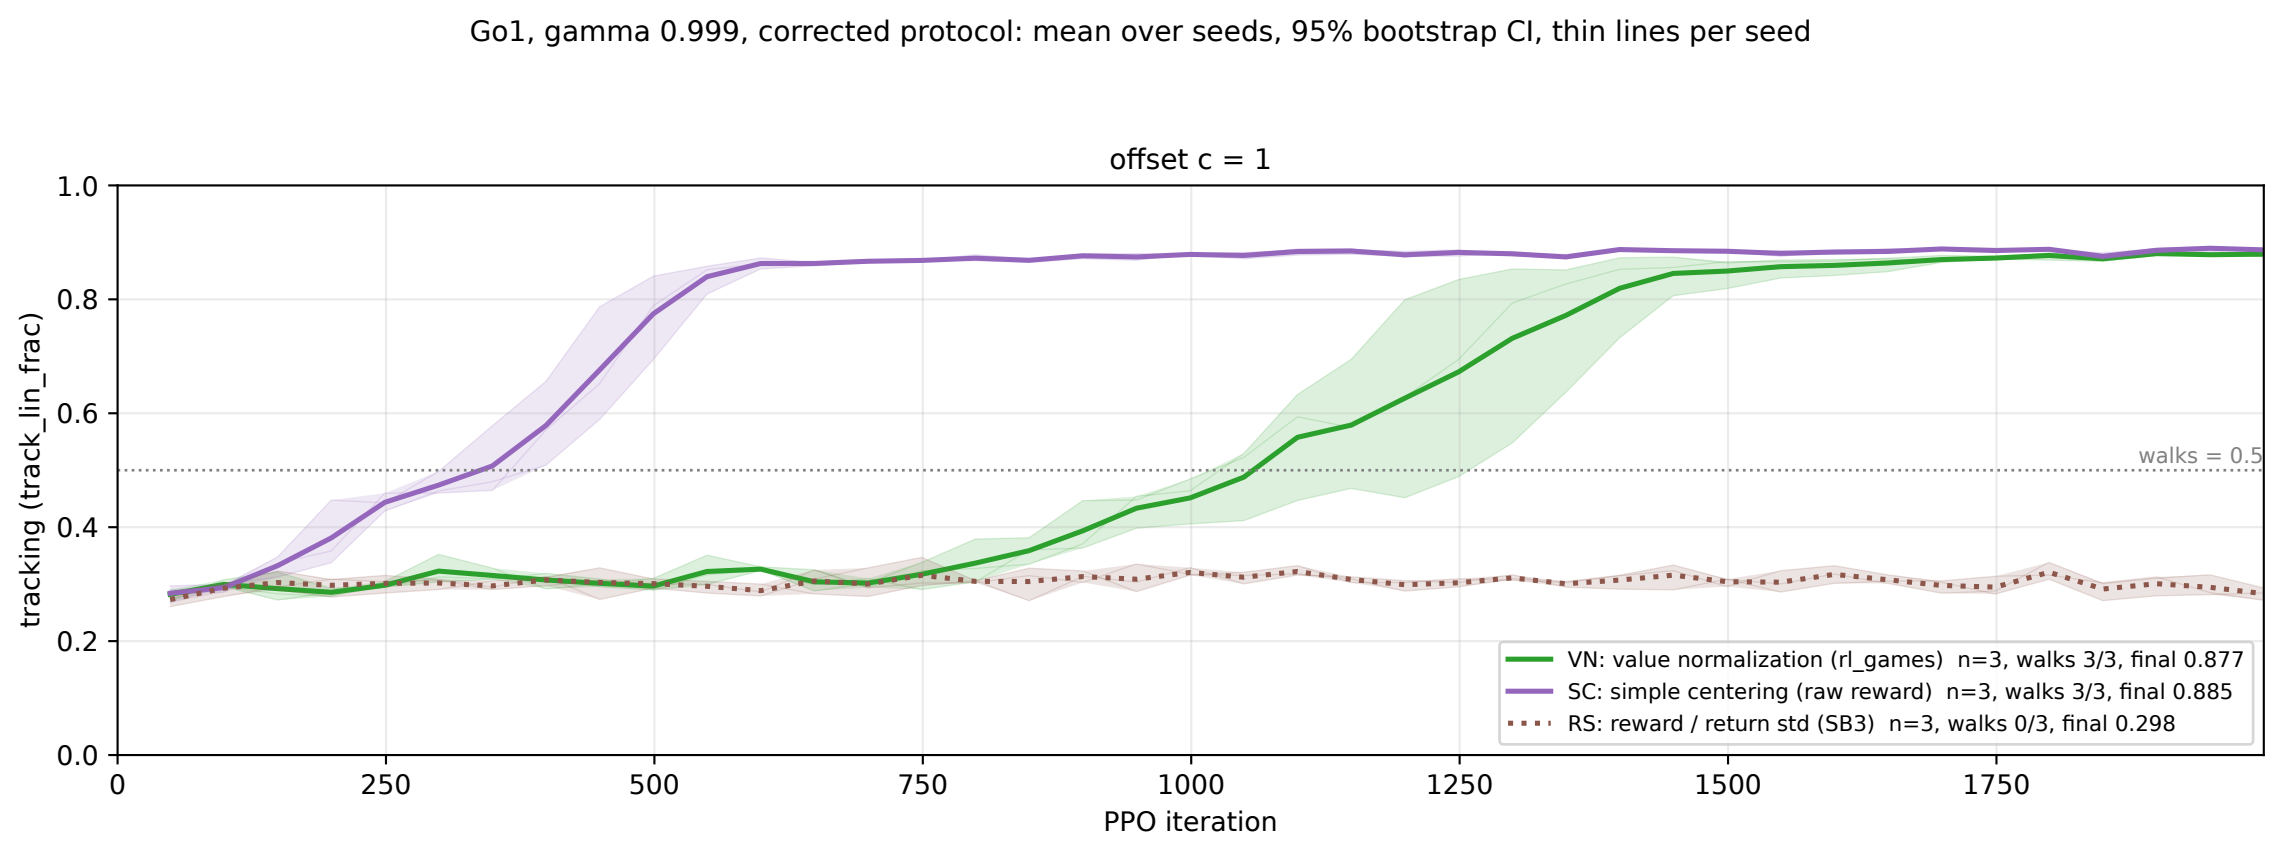
Go1, gamma 0.999, corrected protocol: mean over seeds, 95% bootstrap CI, thin lines per seed
offset c = 1
tracking (track_lin_frac)
PPO iteration
walks = 0.5
VN: value normalization (rl_games) n=3, walks 3/3, final 0.877
SC: simple centering (raw reward) n=3, walks 3/3, final 0.885
RS: reward / return std (SB3) n=3, walks 0/3, final 0.298